

Multimodal NLP & Vision Quiz

Instructions: Select the best answer(s). Multi-select where noted.

Q1. Visual grounding refers to:

- (A) Mapping textual phrases to image regions
- (B) Generating image captions
- (C) Translating images to text
- (D) Segmenting images at pixel level

Q2. (Multi-select) The CLIP model uses which training paradigm?

- (A) Contrastive learning between image and text
- (B) Cross-entropy language modeling
- (C) Self-supervised masked image modeling
- (D) Reinforcement learning

Q3. The multimodal attention mechanism in ViLT fuses modalities by:

- (A) Concatenating visual patches and word tokens into a single Transformer stream
- (B) Training separate encoders with late fusion
- (C) Using a CNN backbone exclusively for images
- (D) Adding textual queries to frozen image features

Q4. In encoder–decoder image captioning, the decoder is typically:

- (A) An RNN or Transformer language model generating text
- (B) A CNN classification head
- (C) An object detector
- (D) A feature extractor only

Q5. (Multi-select) Challenges in multimodal QA include:

- (A) Modality alignment
- (B) Table normalization
- (C) Reasoning across modalities
- (D) High memory consumption

Q6. The term *implicit decomposition* in multimodal QA indicates:

- (A) Splitting complex questions into sub-queries internally
- (B) Using retrieval-based pipelining
- (C) Fine-tuning only the final layer
- (D) Enforcing symbolic reasoning graphs

Q7. A common metric for image captioning that measures n-gram overlap is:

- (A) BLEU
- (B) FID
- (C) WER
- (D) AUC