Reinforcement Learning Quiz

Which algorithms belong to the policy-gradient family?
A) Proximal Policy Optimization (PPO)
B) Q-learning
C) REINFORCE
D) Deep Q-Network (DQN)

An environment is considered partially observable when:
A) The agent can observe the full state
B) Observations lack some hidden variables
C) The reward signal is stochastic
D) Transition dynamics are deterministic

In Q-learning, the key difference from SARSA is:
A) On-policy vs. off-policy nature
B) Use of neural networks
C) Continuous action space support
D) Model-based planning

The Advantage term in Actor–Critic methods is calculated as:
A) State-value minus action-value
B) Action-value minus state-value
C) Reward minus entropy
D) Policy-gradient estimate

Techniques to stabilize deep-RL training include:
A) Experience replay
B) Target networks
C) Layer normalization
D) Early stopping

A discount factor γ close to 0 emphasizes:
A) Immediate rewards
B) Long-term rewards
C) Exploration
D) Deterministic policies

The exploration–exploitation dilemma refers to:
A) Balancing policy and value networks
B) Choosing between trying new actions and using known good actions
C) Data augmentation
D) Hyper-parameter tuning